

SCC : rckpc

2011.11.07

이 용 준

RCKPC COMMUNICATION

Breakdown of SCC LUT Entries

Space	Size	Nbr of LUT Entries
Private Space (system physical memory)	672MB	42 LUT Entries
Shared Space (system physical memory)	64MB	4 LUT Entries
MPB Space (on-die memory)	384MB	24 LUT Entries
SysCGH Space (on-die memory)	384MB	24 LUT Entries
VRC Space	Addresses the VRC configuration registers	1 LUT Entry
MCPC TCP/IP	Addresses the System Interface; access is handled by the PCIe driver	1 LUT Entry

Table 15 Breakdown of LUT Entries for 32GB of System Memory

System Interface

- **System Interface**
 - The System Interface (MCPC TCP/IP) : 16MB
 - Addressable by the core

LUT #	Physical Address	
255	FFFFFFFF - FF000000	Private
254	FEFFFFFF - FE000000	
253	FDFFFFFF - FD000000	
252	FCFFFFFF - FC000000	
251	FBFFFFFF - FB000000	VRC
250	FAFFFFFF - FA000000	Management Console TCP/IP Interface
249	F9FFFFFF - F9000000	
248	F8FFFFFF - F8000000	
247	F7FFFFFF - F7000000	System Configuration Register -- Tile 23

crb & Rckpc

- **Communication between MCPC(the host) & SCC cores**
 - MCPC : crb
 - mcpc_driver/crbif_net.c
 - A network interface for transferring data with SCC cores
 - SCC cores : Rckpc
 - Linuxkernel/linux-2.6.16-mcemu/drivers/net/rckpc.c
 - A layer for transferring data to/from MCPC

Rckpc

- **From Rckpc's point of view**

- `MAX_PACKET_SIZE` : 4096
 - Maximum size of data transfers
- `rxPacketSlots` : 16
 - 1 slot : `MAX_PACKET_SIZE` = 4096
- `rxBaseAddress` : 0x80200000

:	:	:
132	84FFFFFF - 84000000	
131	83FFFFFF - 83000000	Shared MCH3 - 8MB
130	82FFFFFF - 82000000	Shared MCH2 - 8MB
129	81FFFFFF - 81000000	Shared MCH1 - 8MB
128	80FFFFFF - 80000000	Shared MCH0 - 8MB
127	7FFFFFFF - 7F000000	
:	:	:

- Start address of the network buffer in the shared memory space
- 4 consecutive memory tiles are reserved for the 4 quadrants
- Each core uses `rxPacketSlots*MAX_PACKET_SIZE` bytes
- `txMailbox` : 0xFA000000

250	FAFFFFFF - FA000000	Management Console TCP/IP Interface
-----	---------------------	-------------------------------------

- Address of the Tx mailbox

Rckpc

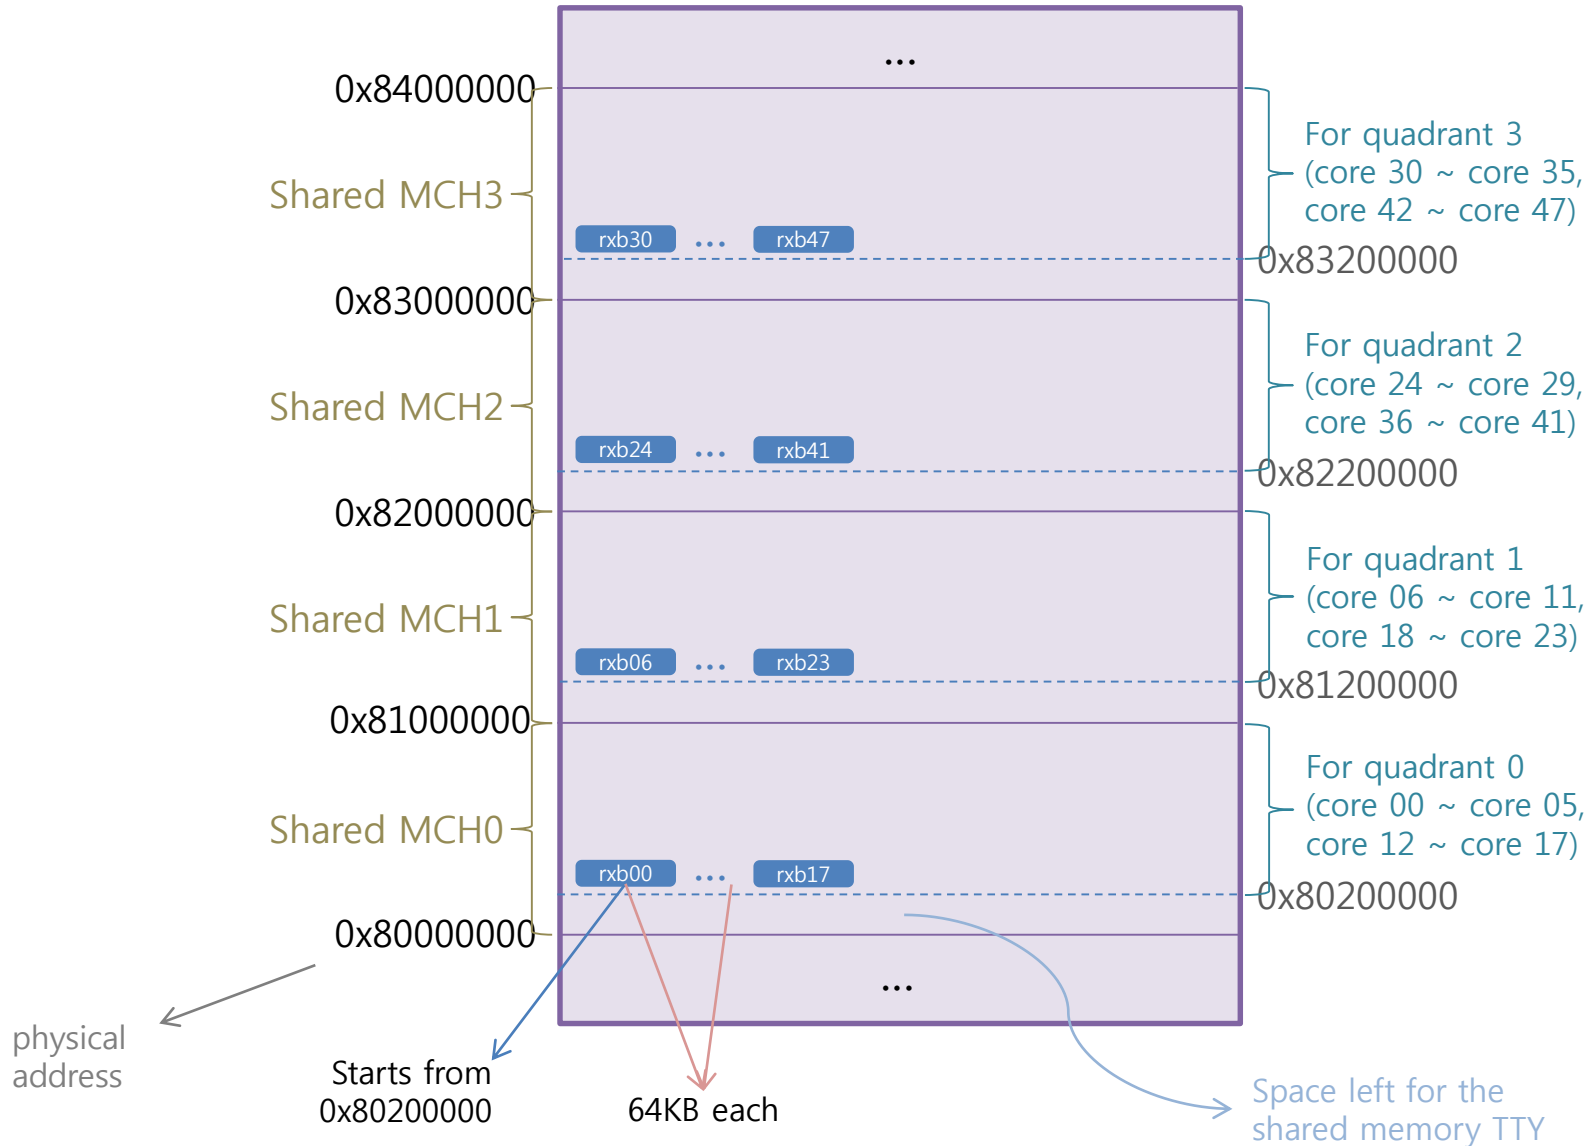
- **struct rckpc_priv @ rckpc.c**

- struct rckpc_priv

```
struct rckpc_priv {  
    ...  
    void* rxb;           // Receive Buffer  
    void* mailbox;       // Host mailbox  
};
```

- rxb
 - Receive buffer
 - Mapped space for rxPacketSlots
 - $\text{rxPacketSlots} * \text{MAX_PACKET_SIZE} = 16 * 4\text{K} = 64\text{K}$ per core
 - Marked as MPB so it can be easily invalidated
 - Write-through flag is also set
 - So every write is immediately propagated to main memory
- mailbox
 - Mapped for txMailbox
 - Offsets of the host mailbox registers
 - MBX_NULL : 0x00
 - MBX_CONFIG : 0x20
 - MBX_PACKETSTART : 0x40
 - MBX_PACKETDATA : 0x60
 - MBX_RXDONE : 0x80

Utilization of Shared Memory



SCC Core Rckpc Send

- **SCC core Rckpc send (`rckpc_tx()`)**
 - Copy the packet data to the mailbox, which is the system IF.
 - If the data is smaller than `RCK_CLINE_SIZE(32B)`,
 - Copy the first cacheline to `MBX_PACKETSTART`.
 - Copy all remaining data to `MBX_PACKETDATA`.
 - If the remaining data is smaller than `RCK_CLINE_SIZE`,
 - Copy `RCK_CLINE_SIZE` bytes to `MBX_PACKETDATA`.
 - And repeat this process.
 - Else,
 - Just copy the remaining bytes to `MBX_PACKETDATA`.
 - Else,
 - Copy them all to `MBX_PACKETSTART`.
 - The source core ID is at the packet header.

SCC Core Rckpc Receive

- **SCC core Rckpc receive (`rckpc_rx()`)**
 - 1. Interrupt triggers the receiver's polling function.
 - 2. The polling function looks for valid packets.
 - The address of the receiving packet is calculated.
 - The current slot at the receiving core's `rxb`
 - Checks for valid packet lengths there.
 - 3. Copy data from the address.
 - 4. Invalidate the packet length so as not to process the data again.
 - 5. Tell the host we are done with the slot and move on to the next.
 - Write the current slot value to `MBX_RXDONE` in the mailbox.
 - Update the current slot value.

crb

- **From crb's point of view**
 - `MAX_PKT_SIZE : 4096`
 - Maximum size of data transfers to Rckpc
 - `txPacketSlots : 16`
 - 1 slot : `MAX_PKT_SIZE = 4096`
 - Same as `rxPacketSlots` in Rckpc code.
 - `txBaseAddress : 0xF0200000`
 - Local address where the shared memory is located
 - At the destination DDR3 memory controller
 - 16 MB at `0xF0000000` of each MC are mapped to `0x80000000 – 0x83000000`.
 - Start address of the packet space at the MC
 - Corresponds to `rxBaseAddress` in Rcp code.

MCPC crb Send

- **MCPC crb send (crbnet_tx())**
 - 1. Extract the destination address from the IP header.
 - Get the core ID.
 - 2. Calculate the address of the buffer(rxb) for the destination.
 - 3. Write the packet data to the address.
 - 32 bytes at a time
 - 4. Generate an interrupt for the destination core.

MCPC crb Receive

- **MCPC crb receive**

- The crbif daemon fetches data from the FPGA.
- Then the filter function passes the packets that are intended for the network interface to crb.
 - `crbif_main.c : crbif_filter()`
 - 1. Decode the packet header.
 - Gets the source core ID.
 - 2. Calls `crbnet_pktHandler()`.

MCPC crb Receive

- **MCPC crb receive**
 - `crbif_net.c : crbnet_pktHandler()`
 - Core ID of the source already known at this stage
 - Decode the different mailbox addresses.
 - If the address is at `MBX_PACKETSTART`,
 - Call `crbnet_tx_start()`.
 - Copy the data from the mailbox.
 - If the address is at `MBX_PACKETDATA`,
 - Call `crbnet_rx_data()`.
 - Copy the data.
 - If the address is at `MBX_RXDONE`,
 - Call `crbnet_rx_complete()`.
 - This is a verification that a send from here to a core's Rckpc has been received fine.

ISSUES WITH RCKPC

Migration & Rckpc

- **Rckpc send & crb receive**
 - After migration,
 - At send,
 - The core IP info is in the packet.
 - **MODIFY : Modify IP to that of the migration src core.**
 - Rckpc shall simply write this packet to the system IF.
 - At receive,
 - crb shall decode the packet header to find out the src IP.
 - **Should be fine as it is.**
 - To solve this problem,
 - **Every core has to maintain a IP-to-Core table.**

Migration & Rckpc

- **crb send & Rckpc receive**
 - After migration,
 - At send,
 - crb shall write at the original rxb of the migrated core.
 - Then, crb shall send an interrupt to the wrong core.
 - **MODIFY : Send the interrupt to the migration dst core.**
 - At receive,
 - When the core gets interrupted,
 - Rckpc shall probably read the original rxb .
 - **Should be fine as it is.**
 - To solve this problem,
 - **MCPC has to maintain a IP-to-Core table.**
 - **At every migration, MCPC has to be told about it somehow.**